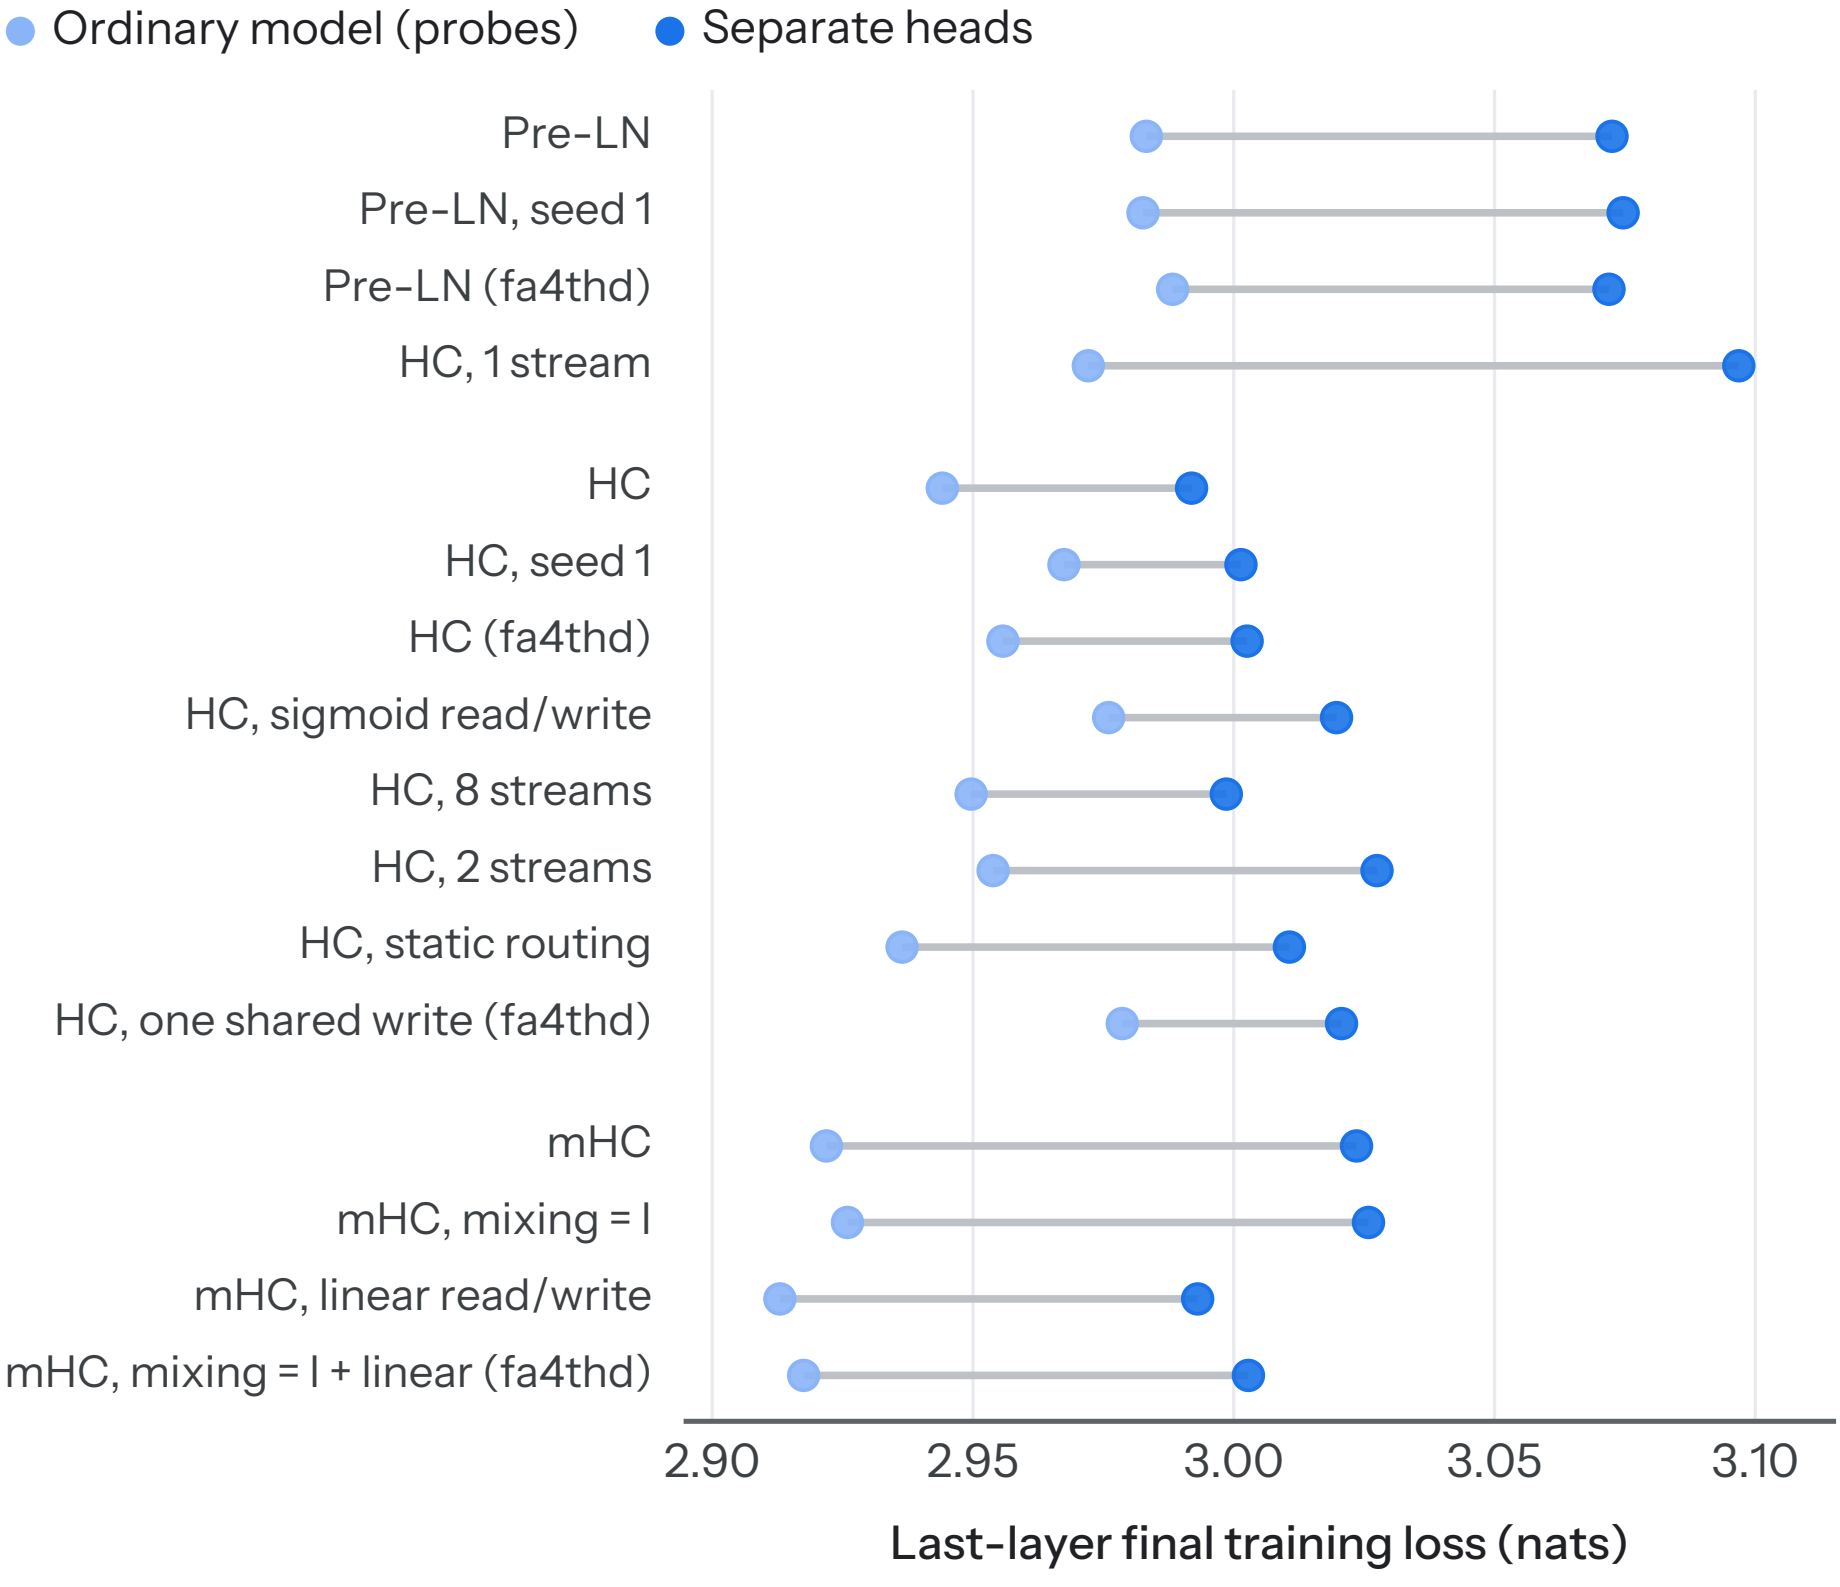
Ordinary model (probes)
Separate heads
Pre-LN
Pre-LN, seed 1
Pre-LN (fa4thd)
HC, 1 stream
HC
HC, seed 1
HC (fa4thd)
HC, sigmoid read/write
HC, 8 streams
HC, 2 streams
HC, static routing
HC, one shared write (fa4thd)
mHC
mHC, mixing = I
mHC, linear read/write
mHC, mixing = I + linear (fa4thd)
2.90
2.95
3.00
3.05
3.10
Last-layer final training loss (nats)